Number of Processes	CPU time over 10000 pixels (ms)
2	6663.93
	6639.5
	6559.31
	6551.18
	6505.81
	6691.15
	6703.51
	6801.81
	6668.67
	6553.03
average	6633.79
4	3552.76
	3529.55
	3500.89
	3561.89
	3468.73
	3625.66
	3645.89
	3420.85
	3451.28
	3512.11
average	3526.961
8	4290.67
	4078.04
	4119.96
	4280.51
	4073.51
	4051.32
	4146.89
	4177.76
	4057.66
	4157.45
average	4143.377
16	5057.54
	5047.86

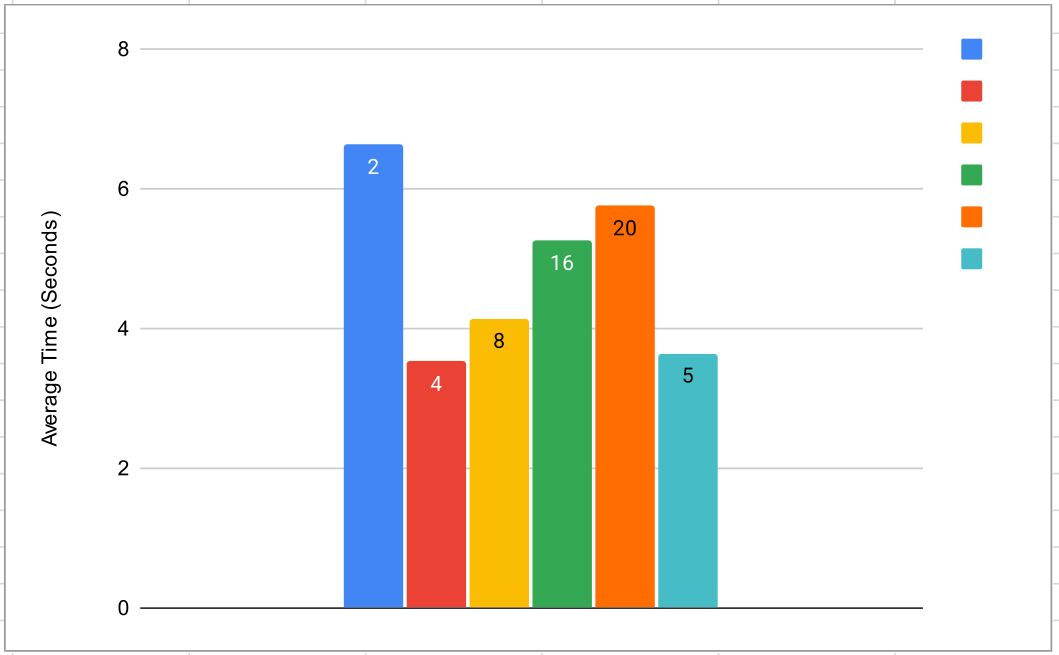


```
Device 0: "NVIDIA GeForce GT 1030"
  CUDA Driver Version / Runtime Version          12.8 / 11.6
  CUDA Capability Major/Minor version number:    6.1
  Total amount of global memory:                 1993 MBytes (2089943040 bytes)
  (003) Multiprocessors, (128) CUDA Cores/MP:    384 CUDA Cores
  GPU Max Clock rate:                            1380 MHz (1.38 GHz)
  Memory Clock rate:                             1050 Mhz
  Memory Bus Width:                              64-bit
  L2 Cache Size:                                 524288 bytes
  Maximum Texture Dimension Size (x,y,z)         1D=(131072), 2D=(131072, 65536), 3D=(16384, 16384, 16384)
  Maximum Layered 1D Texture Size, (num) layers  1D=(32768), 2048 layers
  Maximum Layered 2D Texture Size, (num) layers  2D=(32768, 32768), 2048 layers
  Total amount of constant memory:               65536 bytes
```

	5212.33
	5162.55
	5389.42
	5409.01
	4971.88
	5327.7
	5356.58
	5659.83
average	5259.47
20	5736.61
	5192.16
	5931.57
	5835.46
	5924.04
	5518.8
	5702.09
	5545.28
	6187.97
	5988.88
average	5756.286
5	3597.46
	3570.69
	3688.54
	3629.38
	3763.93
	3573
	3638.72
	3764.86
	3642.28
	3606.56
average	3647.542
	Full image (360000000 pixels):
	4128280 ms

```
  Total amount of shared memory per block:       49152 bytes
  Total shared memory per multiprocessor:        98304 bytes
  Total number of registers available per block: 65536
  Warp size:                                     32
  Maximum number of threads per multiprocessor:  2048
  Maximum number of threads per block:           1024
  Max dimension size of a thread block (x,y,z): (1024, 1024, 64)
  Max dimension size of a grid size    (x,y,z): (2147483647, 65535, 65535)
  Maximum memory pitch:                          2147483647 bytes
  Texture alignment:                             512 bytes
  Concurrent copy and kernel execution:          Yes with 2 copy engine(s)
  Run time limit on kernels:                     No
  Integrated GPU sharing Host Memory:            No
  Support host page-locked memory mapping:       Yes
  Alignment requirement for Surfaces:            Yes
  Device has ECC support:                        Disabled
  Device supports Unified Addressing (UVA):      Yes
  Device supports Managed Memory:                Yes
  Device supports Compute Preemption:            Yes
  Supports Cooperative Kernel Launch:            Yes
  Supports MultiDevice Co-op Kernel Launch:      Yes
  Device PCI Domain ID / Bus ID / location ID:   0 / 134 / 0
  Compute Mode:
     < Default (multiple host threads can use ::cudaSetDevice() with device simultaneously) >

deviceQuery, CUDA Driver = CUDART, CUDA Driver Version = 12.8, CUDA Runtime Version = 11.6, NumDevs = 1
Result = PASS
[u1097883@notchpeak2 ~]$
```