Block Size	GPU time over 8000 pixels
8,8	7278.98
	7209.73
	7129.2
	7289.51
	7141
	7254.09
	7069.13
	7063.89
	7294.01
	7145.09
average	7.187463
16,16	4858.87
	4907.76
	4683.04
	4731.59
	4727.95
	4735.86
	4734.18
	4735.42
	4899.23
average	4.779322222
32,32	27401
	27385
	27411.8
	27418
	27390.4
	27192.9
	27388.4
	27362.6
	27194.2
average	27.34936667
32,16	4787.03

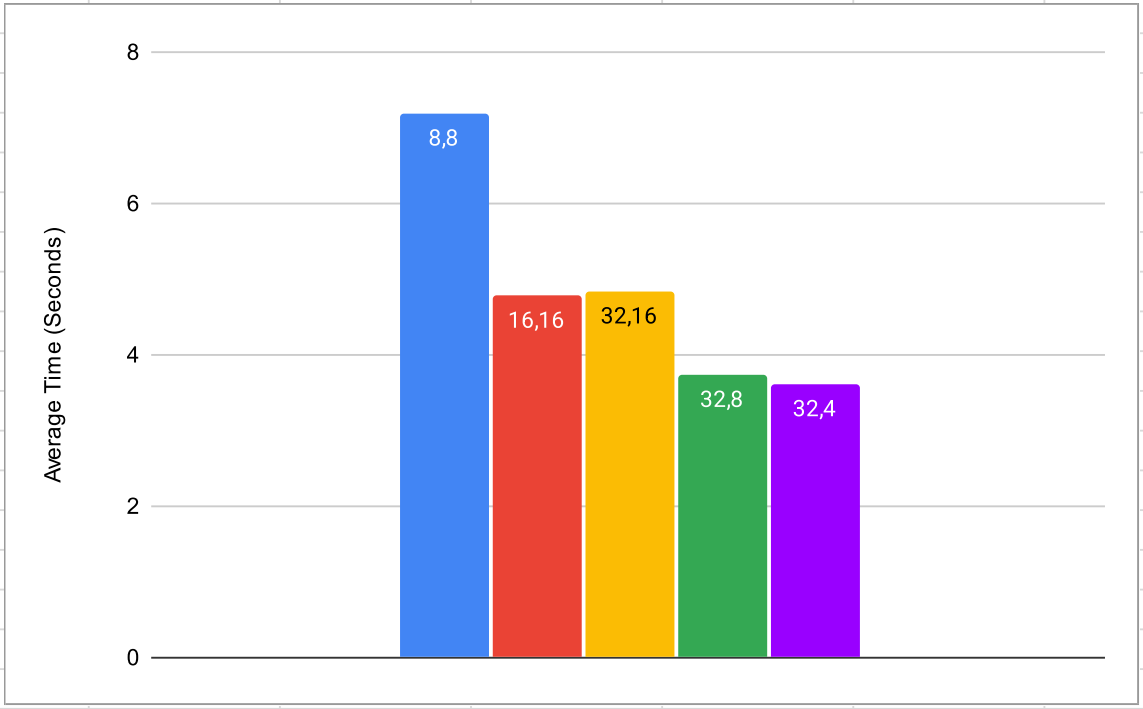


```
Device 0: "NVIDIA GeForce GT 1030"
  CUDA Driver Version / Runtime Version          12.8 / 11.6
  CUDA Capability Major/Minor version number:    6.1
  Total amount of global memory:                 1993 MBytes (2089943040 bytes)
  (003) Multiprocessors, (128) CUDA Cores/MP:    384 CUDA Cores
  GPU Max Clock rate:                            1380 MHz (1.38 GHz)
  Memory Clock rate:                             1050 Mhz
  Memory Bus Width:                              64-bit
  L2 Cache Size:                                 524288 bytes
  Maximum Texture Dimension Size (x,y,z)         1D=(131072), 2D=(131072, 65536), 3D=(16384, 16384, 16384)
  Maximum Layered 1D Texture Size, (num) layers  1D=(32768), 2048 layers
  Maximum Layered 2D Texture Size, (num) layers  2D=(32768, 32768), 2048 layers
```

	4769.92
	4972.55
	4767.16
	4803.19
	4952.36
	4899.55
	4802.22
	4774.48
average	4.836495556
32,8	3873.84
	3804.64
	3708.35
	3702.18
	3702.06
	3702.23
	3702.01
	3701.82
	3701.65
average	3.733197778
32,4	3725.26
	3707.35
	3555.14
	3547.09
	3547.18
	3583.97
	3546.73
	3552.57
	3727.66
average	3.610327778

```
  Total amount of constant memory:               65536 bytes
  Total amount of shared memory per block:       49152 bytes
  Total shared memory per multiprocessor:        98304 bytes
  Total number of registers available per block: 65536
  Warp size:                                     32
  Maximum number of threads per multiprocessor:  2048
  Maximum number of threads per block:           1024
  Max dimension size of a thread block (x,y,z): (1024, 1024, 64)
  Max dimension size of a grid size    (x,y,z): (2147483647, 65535, 65535)
  Maximum memory pitch:                          2147483647 bytes
  Texture alignment:                             512 bytes
  Concurrent copy and kernel execution:          Yes with 2 copy engine(s)
  Run time limit on kernels:                     No
  Integrated GPU sharing Host Memory:            No
  Support host page-locked memory mapping:       Yes
  Alignment requirement for Surfaces:            Yes
  Device has ECC support:                        Disabled
  Device supports Unified Addressing (UVA):      Yes
  Device supports Managed Memory:                Yes
  Device supports Compute Preemption:            Yes
  Supports Cooperative Kernel Launch:            Yes
  Supports MultiDevice Co-op Kernel Launch:      Yes
  Device PCI Domain ID / Bus ID / location ID:   0 / 134 / 0
  Compute Mode:
     < Default (multiple host threads can use ::cudaSetDevice() with device simultaneously) >

deviceQuery, CUDA Driver = CUDART, CUDA Driver Version = 12.8, CUDA Runtime Version = 11.6, NumDevs = 1
Result = PASS
```